

## Fundamentals of Neural Network Implementation

### Introduction

Neural networks are an important part of modern artificial intelligence and machine learning, inspired by the functioning of the human brain. They are designed to recognize patterns, approximate functions, and make predictions based on data. Implementing a neural network requires a deep understanding of its fundamental components, training methodologies, optimization techniques, and architectural considerations. This documentation provides a comprehensive overview of these core aspects.

### Neural Network Structure

A neural network consists of interconnected layers of neurons, each performing a weighted computation on the input data. The primary types of layers include:

1. **Input Layer** – This layer receives raw data and passes it to the next layer without any computation.
2. **Hidden Layers** – These layers perform weighted computations, apply activation functions, and learn intermediate representations from the input data.
3. **Output Layer** – The final layer produces the network's prediction or classification results based on learned features.

Each neuron in a hidden or output layer computes a weighted sum of its inputs, applies a bias, and processes the result using an activation function. The output is then transmitted to the next layer.

### Activation Functions

Activation functions introduce non-linearity, enabling neural networks to model complex relationships in data. Common activation functions include:

- **Sigmoid**: Maps input values to a range between 0 and 1, often used for binary classification.
- **Tanh**: Similar to sigmoid but outputs values between -1 and 1.

- **ReLU (Rectified Linear Unit):** Outputs zero for negative inputs and the input itself for positive values.
- **Softmax:** Converts logits into probabilities for multi-class classification tasks.

The choice of activation function significantly impacts the training process and model performance.

### Training Process

Neural networks learn by adjusting their internal weights based on training data. The key steps in training include:

1. **Forward Propagation** – Input data flows through the network, with each neuron computing an output based on its weights and activation function.
2. **Loss Calculation** – A loss function evaluates the network's prediction against the ground truth. Common loss functions include mean squared error (MSE) for regression and cross-entropy loss for classification.
3. **Backpropagation** – The network calculates gradients using the chain rule of differentiation, determining how much each weight contributes to the error.
4. **Optimization and Weight Update** – An optimization algorithm (such as gradient descent) updates the weights to minimize the loss function iteratively.

### Optimization Algorithms

Optimization plays an important role in training neural networks. Some of the most widely used optimization techniques include:

- **Gradient Descent:** Updates weights by moving them in the opposite direction of the gradient of the loss function.
- **Stochastic Gradient Descent (SGD):** Updates weights after computing the gradient from a single training example, improving convergence speed.
- **Adam (Adaptive Moment Estimation):** Adjusts learning rates dynamically using first and second moments of gradients, balancing speed and stability.
- **RMSprop:** Uses a moving average of squared gradients to normalize learning rates, preventing large updates.

Selecting an appropriate optimizer depends on the problem at hand and the nature of the dataset.

### Regularization Techniques

Overfitting is a common issue in neural networks, where the model memorizes training data instead of generalizing to new examples. Regularization techniques help mitigate this problem:

- **Dropout:** Randomly drops neurons during training to prevent over-reliance on specific pathways.
- **L1 and L2 Regularization:** Adds penalties to large weights.
- **Batch Normalization:** Normalizes inputs at each layer to stabilize training.
- **Early Stopping:** Halts training when validation performance stops improving to avoid overfitting.

### Implementation Considerations

When implementing a neural network, practical aspects must be considered:

- **Data Preprocessing:** Normalizing, scaling, and augmenting data improves model performance.
- **Hyperparameter Tuning:** Adjusting learning rates, batch sizes, and layer sizes optimizes results.
- **Hardware Considerations:** Training large models often requires GPUs or TPUs to accelerate computations.
- **Evaluation Metrics:** Accuracy, precision, recall, and F1-score help assess model effectiveness.

### PyTorch

PyTorch is an open-source machine learning framework that has gained significant popularity among researchers and practitioners in deep learning. Developed by Facebook's AI Research lab (FAIR), PyTorch provides a flexible, Pythonic interface for building and training neural networks. It is widely used in academia and industry due to its intuitive design, strong GPU acceleration, and dynamic computation graph, which makes debugging and experimentation easier compared to traditional static graph frameworks.

One of PyTorch's most notable features is its ability to define and modify neural networks dynamically, rather than requiring a predefined structure before execution. This approach allows researchers and developers to experiment more freely with model architectures, making it particularly well-suited for deep learning research. PyTorch also provides an efficient automatic differentiation engine, which is essential for training neural networks using backpropagation.

In addition to its core functionalities, PyTorch offers multiple libraries that extend its capabilities. Libraries such as TorchVision for computer vision, TorchText for natural language processing, and TorchAudio for audio-related tasks make PyTorch a comprehensive tool for various machine learning applications. Furthermore, PyTorch integrates well with other deep learning tools and supports model deployment through TorchScript and ONNX, enabling efficient inference on different platforms.

### Example of Neural Network Implementation

We exemplify the PyTorch implementation of a neural network which solves a regression problem. The training data consists  $[X, y]$  input-output pairs:

- Input  $X$  consists in multiple features: each instance  $X$  contains features  $[x_1, x_2, \dots, x_n]$ . That is,  $X \in \mathbb{R}^n$
- Output  $y$  consists in a single real value, that is,  $y \in \mathbb{R}$ .

The objective is to train a neural network to learn the mapping between inputs  $X$  and outputs  $y$ , so as to be able to reliably predict the corresponding output  $y$  for any input  $X$ .

Thus, we define and train a simple neural network for regression tasks, whose key components are:

1. **Model Definition (RegressionNN):**
  - Fully connected layers with ReLU activations.
2. **Synthetic Dataset Generation (generate\_dataset):**
  - Linear data with Gaussian noise.
3. **Training Process:**
  - Forward pass, loss computation, backpropagation, and weight updates.

### Imports

```
import torch
import torch.nn as nn
```

```
import torch.optim as optim
```

- `torch`: The PyTorch library for deep learning.
- `torch.nn`: Contains modules for defining neural networks.
- `torch.optim`: Provides optimization algorithms such as Adam and SGD.

### Defining the regression neural network

```
class RegressionNN(nn.Module):  
    def __init__(self, input_size, hidden_size1=64, hidden_size2=32):  
        super(RegressionNN, self).__init__()  
        self.fc1 = nn.Linear(input_size, hidden_size1)  
        self.fc2 = nn.Linear(hidden_size1, hidden_size2)  
        self.fc3 = nn.Linear(hidden_size2, 1)  
        self.relu = nn.ReLU()
```

- The `RegressionNN` class inherits from `nn.Module`, the base class for all neural networks in PyTorch.
- The network consists of three fully connected (`Linear`) layers:
  - First layer: `input_size`  $\rightarrow$  `hidden_size1` (default 64 neurons)
  - Second layer: `hidden_size1`  $\rightarrow$  `hidden_size2` (default 32 neurons)
  - Third layer: `hidden_size2`  $\rightarrow$  output (1 neuron)
- The activation function used is **ReLU** (Rectified Linear Unit), which introduces non-linearity to the network.

### Forward propagation

```
def forward(self, x):  
    x = self.relu(self.fc1(x))  
    x = self.relu(self.fc2(x))  
    x = self.fc3(x)  
    return x
```

The forward pass computes:

- `fc1(x)`  $\rightarrow$  Apply ReLU activation.
- `fc2(x)`  $\rightarrow$  Apply ReLU activation.
- `fc3(x)`  $\rightarrow$  Output a single numerical value.

### Training the neural network

```
if __name__ == "__main__":  
  
    # generate dataset  
    num_features = 3  
    num_instances = 100  
    X, y = generate_dataset(num_features, num_instances)
```

- first the data set is generated:
- `num_features = 3`: Each input sample has 3 features.
- `num_instances = 100`: The dataset contains 100 samples.

```
model = RegressionNN(input_size=num_features)  
lossFunc = nn.MSELoss()  
optimizer = optim.Adam(model.parameters(), lr=0.01)
```

- **Model**: An instance of `RegressionNN` with 3 input features.
- **Loss function**: Mean Squared Error (`MSELoss`), suitable for regression tasks.
- **Optimizer**: Adam, an adaptive learning rate optimization algorithm with learning rate `lr=0.01`.

### Training Loop

```
num_epochs = 100  
for epoch in range(num_epochs):  
    optimizer.zero_grad()  
    outputs = model(X)  
    loss = lossFunc(outputs, y)  
    loss.backward()  
    optimizer.step()  
  
    if (epoch+1) % 10 == 0:  
        print(f'Epoch {epoch+1}, Loss: {loss.item():.4f}')
```

- The training runs for **100 epochs**.
- At each epoch:
  1. `optimizer.zero_grad()`: Clears previous gradients.

## Deep Learning – Laboratory 1

2. `model(x)`: Performs a forward pass to compute predictions.
  3. `lossFunc(outputs, y)`: Computes **Mean Squared Error (MSE)** between predictions and ground truth.
  4. `loss.backward()`: Computes gradients via **backpropagation**.
  5. `optimizer.step()`: Updates weights using the computed gradients.
- Every 10 epochs, the loss value is printed to monitor training progress.

### Tasks

1. Split the data into training and test sets. Evaluate the network on the training and test sets - compute the MSE and  $R^2$  (coefficient of determination).
2. Modify the Dataset:
  - Increase the number of features and observe how it affects the training.
  - Increase the amount of noise (`0.1 * torch.randn(...)`) and analyze its impact on the model's performance.
3. Change the Activation Function
  - Replace the ReLU activation function with LeakyReLU, Sigmoid, or Tanh.
  - Compare the model's performance with different activations.
4. Modify the Loss Function
  - Replace Mean Squared Error (MSE) with Mean Absolute Error (L1Loss).
  - Train the model and compare the effect on loss values.
5. Change the Optimizer
  - Use SGD, RMSprop, or AdaGrad instead of Adam.
  - Compare training speed and convergence behavior.
6. Change the neural network architecture:
  - remove one of the hidden layers
  - add an additional hidden layer

Compare the performance of the networks to the original one.

7. Implement an “early stopping” convergence criterion. That is, stop the training when the loss hasn’t changed significantly over a specified number of epochs (example: if the loss stays within  $\pm 10^{-2}$  during 5 consecutive epochs, stop the training, assuming that no significant improvement will occur over any subsequent epochs).

8. Modify the code so that it implements a binary classification problem, instead of a regression problem.

- the outputs should be labels, instead of real values
- use a loss function more suited to classification (such as Binary Cross Entropy)
- the network should be evaluated in terms of classification accuracy